

R-Zero: SELF-EVOLVING REASONING LLM FROM ZERO DATA

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
Anonymous authors
Paper under double-blind review

ABSTRACT

Self-evolving Large Language Models (LLMs) offer a scalable path toward superintelligence by autonomously generating, refining, and learning from their own experiences. However, existing methods for training such models still rely heavily on vast human curated tasks and labels, typically via fine-tuning or reinforcement learning, which poses a fundamental bottleneck to advancing AI systems toward capabilities beyond human intelligence. To overcome this limitation, we introduce *R-Zero*, a fully autonomous framework that generates its own training data from scratch. Starting from a single base LLM, *R-Zero* initializes two independent models with distinct roles – a **Challenger** and a **Solver**. These models are optimized **separately** and **co-evolve** through interaction: the Challenger is rewarded for proposing tasks near the edge of the Solver’s capability, and the Solver is rewarded for solving increasingly challenging tasks posed by the Challenger. This process yields a targeted, self-improving curriculum without any pre-existing tasks and labels. Empirically, *R-Zero* substantially improves reasoning capability across different backbone LLMs, e.g., boosting the Qwen3-4B-Base by +6.49 on math reasoning benchmarks, and +7.54 on general-domain reasoning benchmarks.

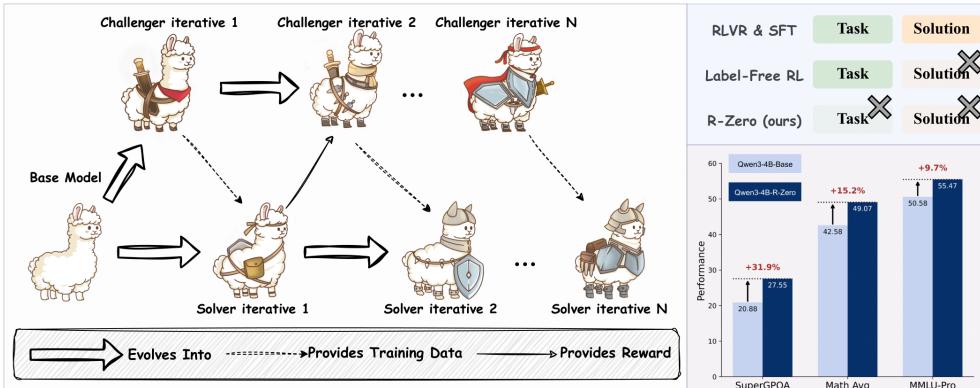


Figure 1: (Left): **R-Zero** employs a co-evolutionary loop between Challenger and Solver. (Right): **R-Zero** achieves strong benchmark gains without any pre-existing tasks or human labels.

1 INTRODUCTION

Self-evolving Large Language Models (LLMs) represent a promising frontier for advancing language intelligence. By autonomously generating, refining, and learning from their own experiences, these models provide a scalable pathway toward artificial superintelligence (Tao et al., 2024; Tan et al., 2024). A critical requirement for training such self-evolving LLMs is access to large volumes of expertly curated tasks and labels, which serve as supervision signals for fine-tuning or reinforcement learning with verifiable rewards (RLVR) (Shao et al., 2024; DeepSeek-AI et al., 2025). However, relying on human annotators to create these tasks and labels is not only costly, labor-intensive, and difficult to scale, but also presents a fundamental bottleneck to advancing AI toward capabilities that could eventually surpass human intelligence (Su et al., 2025; Zhao et al., 2025a).

To reduce dependence on human-curated data, self-generated and label-free methods have been proposed to eliminate the need for explicit supervision. In particular, label-free RL derives reward signals directly from the model’s own outputs, such as sequence-level confidence scores (Li et al., 2025a; Prabhudesai et al., 2025; Huang et al., 2025) and output entropy (Agarwal et al., 2025; Cheng et al., 2025). However, despite removing the need for explicit labels, label-free methods still relies on a pre-existing corpus of tasks, which limits its scalability in truly self-evolving settings. On the other side, self-challenging approaches train LLMs on tasks generated by the models themselves (Zhou et al., 2025a; Wang et al., 2025a; Zhao et al., 2025a). While promising, many of these methods rely on external code executors to ensure that the synthesized tasks are both feasible and verifiable. However, in domains that lack an explicit verification oracle, such as open-ended reasoning, ensuring the quality and correctness of self-generated data remains a significant challenge.

In this paper, we propose *R-Zero*, a framework for training reasoning LLMs that can self-evolve from zero external data. In *R-Zero*, a single base model is initialized with two roles – a **Challenger** and a **Solver** that are independently optimized but **co-evolve** throughout the RL process. During co-evolving, the Challenger is rewarded for generating tasks targeted to be at the edge of Solver’s current abilities, while the Solver is rewarded for solving increasingly challenging tasks posed by the Challenger. Framework details are provided in Section 2, but briefly, in the Challenger training phase, the Challenger is trained via Group Relative Policy Optimization (GRPO) (Shao et al., 2024) to generate difficult questions. The reward signal is derived from the uncertainty for the frozen Solver, which is measured by the self-consistency of its multiple generated answers. In the Solver training phase, the Solver is fine-tuned with GRPO on a filtered set of these challenging questions generated by the now-frozen Challenger, using the pseudo-labels voted by itself. This entire process repeats, creating a self-evolving cycle that operates without any human intervention.

Our experiments demonstrate that *R-Zero* is a model-agnostic framework, consistently and iteratively improving the reasoning abilities of different backbone LLMs. For example, Qwen3-4B-Base model’s average score on math benchmarks increased by a significant **+6.49** points after three iterations of self-evolution. Moreover, the reasoning skills learned through our math-focused questions can generalize to complex general-domain tasks, with models trained using *R-Zero* showing significant improvements on general domain reasoning benchmarks like MMLU-Pro (Wang et al., 2024) and SuperGPQA (Du et al., 2025). Our further analysis finds that *R-Zero* can act as a mid-training method, as models first improved by our method achieve higher performance after fine-tuned on labeled data. In addition, we provide an in-depth analysis that validates our framework’s components, demonstrates its synergy with supervised fine-tuning, and characterizes the co-evolutionary dynamics to identify both strengths and limitations, offering insights for future research.

2 METHOD

2.1 OVERVIEW

We propose ***R-Zero***, a fully automated framework featuring a **Challenger** and a **Solver**, both initialized from the same base LLM. The framework operates in an iterative loop. We illustrate the main framework in Figure 2. First, the Challenger (Q_θ) is trained with Group Relative Policy Optimization (GRPO) to generate synthetic questions that are challenging for the current Solver (Sec. 2.2). A training dataset of question-answer pairs is then constructed from these synthetic questions using a filtering strategy and a majority-vote mechanism (Sec. 2.3). Next, the Solver (S_ϕ) is fine-tuned on this new dataset, also using GRPO (Sec. 2.4). This iterative process allows the Challenger and Solver to co-evolve, leading to a progressively more capable Solver. The entire framework is self-supervised, requiring no human intervention.

2.2 CHALLENGER TRAINING

The Challenger, Q_θ , is an autoregressive language model trained to generate challenging questions. We train Q_θ using the GRPO algorithm detailed in Sec. I. The core of this process lies in designing a reward function that accurately captures the desired properties of a “good” question. This final scalar reward, r_i , is then used in the GRPO. We focus on generating questions specifically within the domain of mathematics, as it provides a convenient and self-contained setting for our framework; the

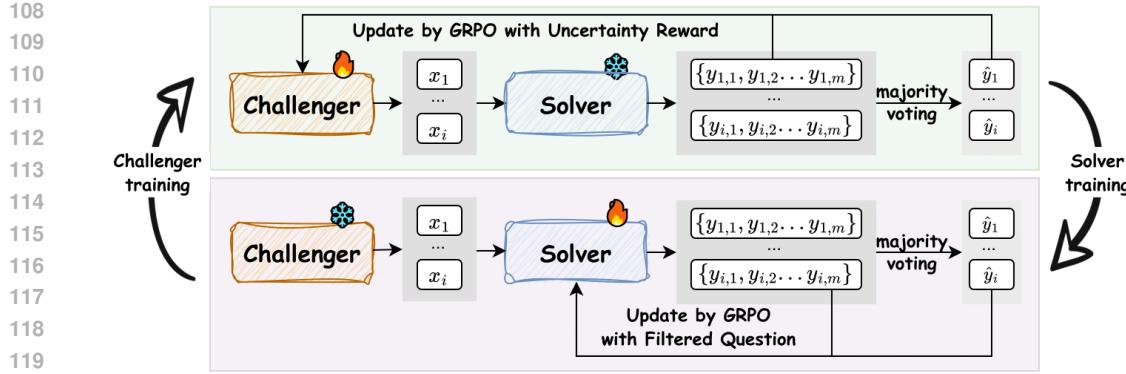


Figure 2: An overview of our *R-Zero* framework, which illustrates the co-evolution of the Challenger and the Solver. **Top:** In the Challenger training phase, the Challenger is trained via GRPO to generate difficult questions. The reward signal is derived from the uncertainty for the frozen Solver, which is measured by the self-consistency of its multiple generated answers. **Bottom:** In the Solver training phase, the Solver is fine-tuned with GRPO on a filtered set of these challenging questions generated by the now-frozen Challenger, using the pseudo-labels voted by itself.

objective nature of mathematical answers allows for the straightforward generation of pseudo-labels via majority voting, without the need for external verification environments like code executors.

Uncertainty Reward. To guide the Challenger toward producing challenging yet solvable questions, we first define an uncertainty score. For a generated question x , we query the current Solver S_ϕ for m responses $\{y_1, \dots, y_m\}$. The most frequent response is treated as the pseudo-label $\tilde{y}(x)$, and we compute the Solver’s empirical accuracy as $\hat{p}(x; S_\phi) = \frac{1}{m} \sum_{j=1}^m \mathbb{1}\{y_j = \tilde{y}(x)\}$. The uncertainty reward is then defined as:

$$r_{\text{uncertainty}}(x; \phi) = 1 - 2 \left| \hat{p}(x; S_\phi) - \frac{1}{2} \right|$$

This function incentivizes questions where the Solver is maximally uncertain (accuracy approaches 50%). We provide a theoretical motivation for this reward function in Appendix F.

Repetition Penalty. To encourage diversity within a training batch \mathcal{X} , we introduce a repetition penalty. We could use any similarity metric, but in our case, we specifically use the BLEU score for faster computation, as this calculation must be performed numerous times during the rollout process. We compute pairwise distances using BLEU score similarity, $d_{ij} = 1 - \text{BLEU}(x_i, x_j)$, and group questions where $d_{ij} < \tau_{\text{BLEU}}$ into clusters $\mathcal{C} = \{C_1, \dots, C_K\}$. The penalty for a question x_i in a cluster C_k is proportional to its relative size:

$$r_{\text{rep}}(x_i) = \lambda \frac{|C_k|}{B}$$

where B is the batch size and λ is a scaling factor. In our experiments, we set $\lambda = 1$. The implementation details are shown in Appendix B.4.

Format Check Penalty. A critical first step in the reward pipeline is a structural format check to verify that each generated question is correctly enclosed within `<question>` and `</question>` tags. If the output does not adhere to this required structure, it is immediately assigned a final reward of 0, and no further reward signals are computed.

Composite Reward and Policy Update. For all questions that pass the format check, we calculate a composite reward. The final scalar reward r_i for each valid question x_i combines signals for uncertainty and repetition:

$$r_i = \max(0, r_{\text{uncertainty}}(x_i; \phi) - r_{\text{rep}}(x_i))$$

With these rewards $\{r_1, \dots, r_G\}$ for a batch of generated questions, we compute the advantage \hat{A}_i for each question and update the Challenger’s policy Q_θ by minimizing the GRPO loss $\mathcal{L}_{\text{GRPO}}(\theta)$.

162 **2.3 SOLVER DATASET CONSTRUCTION**
163

164 After updating the Challenger, we use it to generate a new, curated dataset to train the Solver. This
165 process acts as a curriculum generator. We first sample a large pool of N candidate questions from
166 the Challenger’s policy, $x_i \sim Q_\theta(\cdot | p_0)$. For each question, we obtain m answers from the current
167 Solver, determine the pseudo-label \tilde{y}_i via majority vote, and calculate the empirical correctness \hat{p}_i .
168 A question-answer pair (x_i, \tilde{y}_i) is added to the training set \mathcal{S} only if its correctness falls within an
169 informative band, $|\hat{p}_i - \frac{1}{2}| \leq \delta$. This filtering step discards tasks that are either too easy or too hard.

170 While the primary goal of this filtering is to discard tasks that are too easy or too hard, it also serves
171 as an implicit quality control mechanism. Since our pseudo-labels are derived from a majority vote,
172 a very low empirical correctness \hat{p}_i often indicates that the question itself is ambiguous, ill-posed,
173 or that the resulting pseudo-label is unreliable. By filtering out these low-consistency items, our
174 method simultaneously improves the quality and the uncertainty calibration of the training data.

175 **2.4 SOLVER TRAINING**

177 The Solver, S_ϕ , is then fine-tuned on the curated dataset of challenging problems \mathcal{S} . We also use
178 GRPO for this stage, but with a simpler, verifiable reward signal. For a given question $x_i \in \mathcal{S}$ with
179 its pseudo-label \tilde{y}_i , the Solver generates a batch of answers, each assigned a binary reward r_j :

$$181 \quad r_j = \begin{cases} 1, & \text{if } x_j \text{ matches with the pseudo-label } \tilde{y}_i, \\ 182 \quad 0, & \text{otherwise.} \end{cases}$$

184 This verifiable reward is used to compute the advantage \hat{A}_j , and the Solver’s policy S_ϕ is sub-
185 sequentially updated by minimizing the GRPO loss $\mathcal{L}_{\text{GRPO}}(\phi)$. This process enhances the Solver’s
186 ability to correctly answer the difficult questions generated by its co-evolving Challenger.

187 **3 EXPERIMENTS**

190 **3.1 MODELS AND TRAINING DETAILS**

192 We employ the Qwen3-4B-Base (Yang et al., 2025) and Qwen3-8B-Base models to assess the impact
193 of scale within a single architectural family. Second, to ensure our approach is effective on a dis-
194 tinct lineage, we utilize the OctoThinker-3B and OctoThinker-8B models (Wang et al., 2025b). This
195 choice is particularly relevant as Wang et al. (2025b) reported that applying RL training directly to
196 Llama models yielded suboptimal results. As the OctoThinker series is continually trained from the
197 Llama-3.1 models (Dubey et al., 2024), this comprehensive selection allows us to test our frame-
198 work across different foundational architectures – Qwen vs. Llama. We assess our framework on a
199 comprehensive suite of benchmarks, with the evaluation benchmarks presented in Appendix C.

200 Our entire framework is implemented based on the EasyR1 codebase (Zheng et al., 2025b). In
201 each iteration of the *R-Zero* co-evolutionary loop, we follow a specific set of hyperparameters. The
202 Challenger (Q_θ) first generates a candidate pool of $N = 8,000$ questions. To construct the train-
203 ing dataset for the Solver, these questions are filtered based on consistency. For each candidate
204 question, we sample $m = 10$ answers from the current Solver (S_ϕ). A question is retained for the
205 training set only if the number of answers matching the majority-vote pseudo-label is between 3 and
206 7, inclusive ($\delta = 0.25$). This numerical range is consistent with the methodology used in previous
207 research (Zhang & Zuo, 2025; Li et al., 2025b; Bercovich et al., 2025). When training the Chal-
208 lenger, the uncertainty reward $r(x; \phi)$ is calculated by sampling $m = 10$ responses from the Solver.
209 For the intra-batch repetition penalty, we set the clustering distance threshold to $\tau_{\text{BLEU}} = 0.5$.

210 **3.1.1 EVALUATION SETTING**

211 The evaluation code is adopted from General-Reasoner (Ma et al., 2025). To ensure consistency, we
212 reran the released evaluation code and reported the corresponding results. The reproduced results
213 are well aligned with General-Reasoner and those in the Qwen-3 technical report (Yang et al., 2025).

215 The results presented in all experimental tables are obtained after 45 training steps, while in the
figures we report evaluations performed at checkpoints every 15 steps during solver training. All

Table 1: Comprehensive results on mathematical reasoning benchmarks. We compare each base model against a *R-Zero* (⚡ challenger) baseline (where the Solver is trained on questions from an untrained Challenger) and our method, *R-Zero*. The peak performance is highlighted in **bold**.

Model Name	Avg.	AMC	Minerva	MATH	GSM8K	Olympiad	AIME25	AIME24
<i>Qwen3-4B-Base</i>								
Base Model (w/o training)	42.57	45.70	38.24	68.20	87.79	41.04	10.30	6.7
Absolute Zero	46.42	52.45	41.96	76.20	89.34	42.56	10.20	12.20
<i>R-Zero</i> (⚡ challenger)	45.01	45.00	45.22	72.80	87.87	41.19	10.20	12.80
<i>R-Zero</i> (our method)	49.93	57.27	52.94	79.60	92.12	44.59	9.60	13.40
<i>Qwen3-8B-Base</i>								
Base Model (w/o training)	48.64	51.95	50.00	78.00	89.08	44.74	12.10	14.60
Absolute Zero	52.68	57.89	57.90	76.60	92.00	47.80	18.20	18.40
<i>R-Zero</i> (⚡ challenger)	52.10	60.70	57.72	81.60	92.56	46.44	11.60	14.10
<i>R-Zero</i> (our method)	53.72	61.67	60.66	82.00	94.09	48.89	13.30	15.40
<i>OctoThinker-3B</i>								
Base Model (w/o training)	26.64	17.19	24.26	55.00	73.69	16.15	0.21	0.00
Absolute Zero	27.23	22.50	22.70	53.20	75.80	13.20	0.50	2.70
<i>R-Zero</i> (⚡ challenger)	27.51	20.19	24.63	54.60	74.98	15.70	0.10	2.40
<i>R-Zero</i> (our method)	29.32	27.03	27.57	54.20	74.98	18.22	3.23	0.00
<i>OctoThinker-8B</i>								
Base Model (w/o training)	36.41	32.11	41.91	65.20	86.96	26.52	1.56	0.62
Absolute Zero	36.60	32.50	44.90	62.80	87.00	25.60	3.30	0.10
<i>R-Zero</i> (⚡ challenger)	36.98	29.30	42.28	66.20	88.10	27.56	1.04	4.38
<i>R-Zero</i> (our method)	38.52	34.03	48.22	68.80	87.19	27.56	0.42	3.44

results are reported based on the held-out test sets, ensuring a fair comparison across baselines and reproduced methods. Further implementation details and prompts can be found in Appendix B.

3.2 RESULTS IN MATHEMATICAL REASONING

The comprehensive results of our experiments are presented in Table 1. The findings confirm that our proposed framework, *R-Zero*, is a highly effective, model-agnostic method for enhancing the performance of language models on mathematical tasks across different architectures and scales.

Our iterative training process consistently and substantially improves upon the performance of the base models. This holds true for large models like Qwen3-8B-Base, where three iterations of *R-Zero* raise the average performance from a baseline of 49.18 to 54.69, a significant gain of **+5.51** points. Similarly, on the smaller OctoThinker-3B, our method improves the average score from 26.64 to 29.32 (**+2.68** points), demonstrating the broad applicability of our self-supervised training loop.

The critical role of the Challenger’s RL-based training is validated by the immediate performance leap from the Base Challenger to the first iteration of *R-Zero*. On Qwen3-8B-Base, this first iteration provides a +1.52 point gain over the baseline, and the improvement is even more pronounced on Qwen3-4B-Base at +3.7 points. This confirms that the intelligent curriculum generated by the RL-trained Challenger is significantly more effective than that of a non-trained generator.

3.3 RESULTS IN GENERAL REASONING

Previous work has demonstrated that training language models on reasoning-intensive domains, such as mathematics, can lead to improvements in general-domain capabilities (Huan et al., 2025). A key question, however, is whether this generalization effect still holds when the training curriculum is not human-labeled, but entirely self-generated through *R-Zero*.

As shown in Table 2, this transfer of skills is evident across all tested models. For instance, three iterations of our math-focused training improve the average general-domain score of Qwen3-8B-Base by +5.13 points and OctoThinker-3B by +3.65 points. This generalization also extends to the key performance patterns observed in the mathematical results. This confirms that our method does not merely teach domain-specific knowledge, but enhances the model’s underlying capabilities in a way that successfully generalizes across domains.

270 Table 2: Results on general-domain reasoning benchmarks. The table compares the Base Model, a
 271 *R-Zero* (✿ challenger) baseline, and our *R-Zero*. The peak performance is highlighted in **bold**.
 272

Model Name	Overall Avg.	SuperGPQA	MMLU-Pro	BBEH
<i>Qwen3-4B-Base</i>				
Base Model (w/o training)	26.34	20.88	50.58	7.57
Absolute Zero	29.33	27.10	52.60	8.30
<i>R-Zero</i> (✿ challenger)	28.52	24.77	54.20	6.59
<i>R-Zero</i> (our method)	31.15	27.55	55.47	10.42
<i>Qwen3-8B-Base</i>				
Base Model (w/o training)	31.98	28.33	58.97	8.63
Absolute Zero	34.40	31.89	60.50	10.80
<i>R-Zero</i> (✿ challenger)	31.29	30.12	54.14	9.60
<i>R-Zero</i> (our method)	34.50	31.38	61.53	10.60
<i>OctoThinker-3B</i>				
Base Model (w/o training)	7.47	10.09	10.87	1.46
Absolute Zero	16.03	17.70	24.30	6.10
<i>R-Zero</i> (✿ challenger)	10.04	11.19	14.53	4.40
<i>R-Zero</i> (our method)	11.12	12.44	16.71	4.20
<i>OctoThinker-8B</i>				
Base Model (w/o training)	11.70	13.26	20.21	1.64
Absolute Zero	18.40	18.80	31.40	5.00
<i>R-Zero</i> (✿ challenger)	21.30	16.99	41.46	5.46
<i>R-Zero</i> (our method)	23.00	19.82	40.92	8.25

293 4 ANALYSIS

294
 295 In this section, we conduct a series of in-depth analyses to better understand the behavior and effec-
 296 tiveness of our *R-Zero* framework. To ensure consistency, all analytical experiments presented here
 297 are conducted on the Qwen3-4B-Base model, unless explicitly stated otherwise. Some additional
 298 analyses are shown in Appendix D.
 299

300 4.1 ABLATION STUDY

301  Removing Repetition Penalty and Task Filtering will harm the final performance.

304 To isolate the contribution of each key component within our *R-Zero* framework, we conduct a
 305 comprehensive ablation study on the Qwen3-4B-Base model. We specifically investigate the
 306 importance of two critical modules by disabling them one at a time and observing the impact on
 307 performance. The results are summarized in Table 3.

308 As shown in the table, removing any core
 309 components leads to a significant degradation
 310 in performance. Removing the **Repetition**
 311 **Penalty** harms performance, indicating that
 312 generating a diverse set of questions is crucial
 313 for effective Solver training.

314 Finally, disabling the **Task Filtering** module
 315 results in a notable performance drop, partic-
 316 ularly on the general-domain average, which
 317 falls by over 6 points. As discussed in Sec-
 318 tion 2.3, this filtering serves a dual purpose: it
 319 calibrates the curriculum’s difficulty and acts as an implicit quality control mechanism by removing
 320 questions with low answer consistency. Without this filter, the Solver is trained on noisy data that
 321 likely includes ambiguous or ill-posed questions, which harms its ability to learn robustly.

322 4.2 ITERATION SCALING

323 Table 3: Ablation results on Qwen3-4B-Base. **w/o**
Filtering: Disables the difficulty-based curricu-
 324 lum filtering. **w/o Rep. Penalty**: Removes the
 325 repetition penalty from the Challenger’s reward.

Method	Math AVG	General AVG
<i>R-Zero</i>	49.07	31.15
<i>Ablations</i>		
⊓ w/o Rep. Penalty	45.76	28.73
⊓ w/o Filtering	47.35	26.69

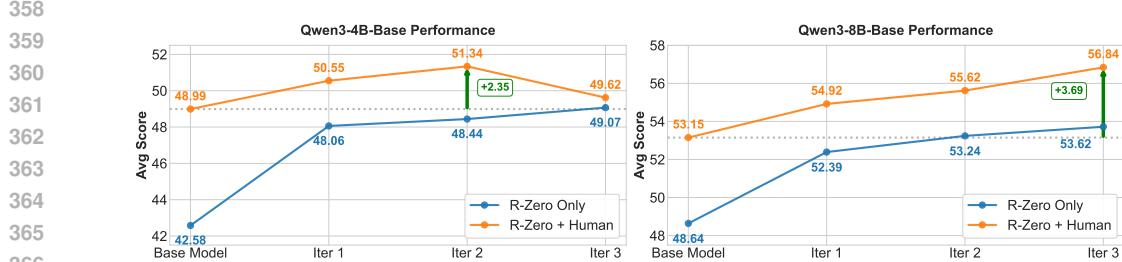
324
 325  Model performance eventually converges, with the timing depending on model size.
 326
 327
 328 Previous results demonstrate that *R-Zero* enhances the Solver’s capabilities. This raises a critical
 329 question about the long-term stability of our self-improvement loop: *what are the limits of this*
 330 *process, and whether the eventual performance degrades?* In this section, we conduct an analysis
 331 to investigate these iteration scaling dynamics, aiming to diagnose the underlying cause of this
 332 instability.

333 As illustrated in Figure 3, our framework initially delivers on its promise, with models of
 334 all sizes showing significant performance improvements in the early stages of co-evolution.
 335 Unfortunately, this virtuous cycle does not continue indefinitely. After multiple iterations, we
 336 observe a consistent and concerning trend of performance degradation across all models. In-
 337 triguely, we found a direct correlation between model scale and resilience to this col-
 338 lypse: the larger the model, the later the onset of performance degradation.
 339

340 For instance, the smallest 0.6B model reaches its peak performance as early as the first iteration
 341 (Step 15), after which its capabilities begin to decline. In contrast, the largest 4B model sustains its
 342 upward trajectory for three full iterations, only experiencing a sharp drop at Step 60. This pattern
 343 strongly suggests that while larger model capacity can delay the negative effects, it does not pre-
 344 vent them. This eventual collapse points to an inherent instability or limitation within our current
 345 self-improvement framework, highlighting a critical area for future investigation. We present some
 346 additional analysis results for this in Appendix E.
 347

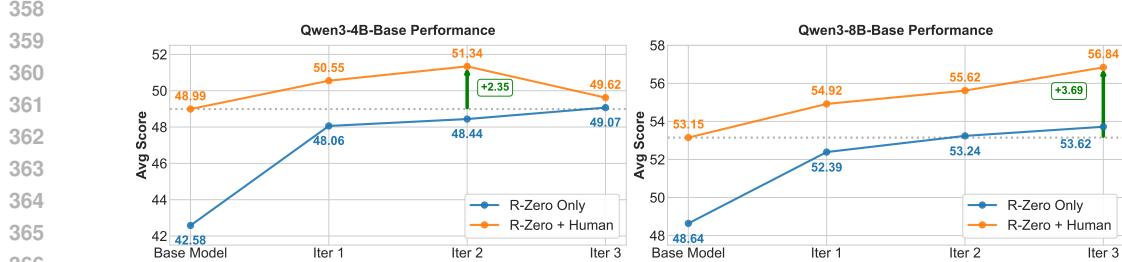
348 4.3 SYNERGY WITH SUPERVISED DATA

349  Using *R-zero* as a mid-training method boosts the effect of later training on human data.
 350
 351 To analyze the utility of our framework in scenarios where a labeled dataset is available, we measure
 352 the synergy between *R-Zero* and traditional supervised fine-tuning using labeled datasets ¹. The
 353 GRPO settings for this experiment were kept identical to our main experiments.



354
 355 Figure 3: Math performance across different iteration times and model scales. The star markers indicate
 356 the peak performance for each model size.
 357

358 To analyze the utility of our framework in scenarios where a labeled dataset is available, we measure
 359 the synergy between *R-Zero* and traditional supervised fine-tuning using labeled datasets ¹. The
 360 GRPO settings for this experiment were kept identical to our main experiments.



361 Figure 4: Performance of *R-Zero* when combined with supervised fine-tuning. The dashed line repre-
 362 sents the baseline of fine-tuning the base model on labelled data alone, showing that our iterative
 363 method provides a better initialization.
 364

365 We first establish a supervised baseline by fine-tuning the base model directly on the labeled data.
 366 For this process, we also employ GRPO.

367 We then apply our *R-Zero* framework, where at the end of each co-evolutionary iteration, the re-
 368 sulting checkpoint is also fine-tuned on the same labeled dataset. The results show that our method
 369 provides significant additional gains. As highlighted in Figure 4, this represents a gain of **+2.35**
 370 **points** over the human-label-only baseline for 4B model.
 371

¹<https://huggingface.co/datasets/hiyouga/math12k>

This finding confirms that *R-Zero* is not redundant with labeled data; instead, it acts as a powerful performance amplifier. The co-evolutionary process enables the model to better leverage the supervised information and achieve performance levels unattainable by standard fine-tuning alone.

We also studied the scenario where human-labeled datasets are mixed into the *R-Zero* generated dataset for training. We found that this concurrent training approach outperforms using either *R-Zero* dataset or humanlabel dataset alone, but it does not surpass the sequential strategy of first applying *R-Zero* and then performing SFT. We hypothesize that mixing the labeled data during *R-Zero* training may dilute the high-quality signal while partially mitigating noise. Performing *R-Zero* first allows the model to acquire reasoning ability before leveraging the high-quality human-labeled data, which appears to be the most effective strategy. The precise underlying reasons require further investigation.

Dataset	AMC	Minerva	MATH	GSM8K	Olympiad	AIME25	AIME24	SuperGPQA	MMLU-Pro	BBEH
Human	57.97	55.15	80.8	92.04	48	9.58	10.31	29.49	57.03	9.71
<i>R-Zero</i>	57.27	52.94	79.6	92.12	44.59	9.6	13.4	27.55	55.47	10.42
<i>R-Zero+Human</i>	57.9	53.2	81.2	92.2	47.7	10.32	14.67	29.4	58.2	11.8

Table 4: Performance comparison across benchmarks when training with human-labeled data, *R-zero* self-generated data, and a mixed training strategy.

4.4 EVOLUTION OF QUESTION DIFFICULTY AND DATA ACCURACY

 Question difficulty increases progressively with decreasing pseudo-label accuracy.

Table 5: Performance and data accuracy analysis. The highlighted column represents the *true accuracy* of the self-generated pseudo-labels for each question set.

Performance of Evaluated Model (vs. Ground Truth)					
	Base Model	Solver (step 15)	Solver (step 30)	Solver (step 45)	Pseudo-Label Acc.
$\mathcal{D}_{\text{Step 15}}$	48.0	59.0	57.0	61.0	79.0%
$\mathcal{D}_{\text{Step 30}}$	52.5	53.0	51.5	53.5	69.0%
$\mathcal{D}_{\text{Step 45}}$	44.0	47.0	45.0	50.5	63.0%

To understand the co-evolutionary dynamic, we analyzed how the Challenger’s generated questions and their corresponding pseudo-labels change across iterations. We sampled 200 questions from the Challenger’s policy after each of the first three training iterations, creating three distinct test sets: $\mathcal{D}_{\text{step 15}}$, $\mathcal{D}_{\text{step 30}}$, and $\mathcal{D}_{\text{step 45}}$. For this analysis, we assumed the external oracle model, GPT-4o, to be a perfect annotator, providing the ground truth answers for all generated questions.

The evaluation was conducted as follows: the performance of our internal models was measured against these GPT-4o ground truth answers. The score reported for GPT-4o itself, however, reflects the **true accuracy of our self-generated pseudo-labels** by comparing the pseudo label against the ground truth from the oracle (GPT-4o). The results on the filtered dataset are summarized in Table 5.

This analysis reveals a multi-faceted dynamic. The first finding is that the questions generated by the Challenger become **progressively more difficult**. This is directly evidenced by evaluating a fixed model against the evolving question sets. For instance, the performance of the static Solver (Step 15), when measured against the consistent GPT-4o ground truth, drops from 59.0% on the questions for the Step 15 training to 47.0% on the questions for the Step 45. This confirms that the Challenger is successfully increasing the intrinsic difficulty of its curriculum.

The second finding, revealed by the highlighted column, pertains to the **true accuracy of the self-generated dataset**. Unfortunately, while the accuracy of the pseudo-labels is initially high at 79.0%, it systematically drops to 63.0% by the third iteration. This trend indicates that as the system generates more difficult problems, the Solver’s majority vote becomes a less reliable source for ground truth. This decline in data quality is a critical trade-off and a potential bottleneck for the framework’s ultimate performance.

432 5 RELATED WORK

434 5.1 LABEL-FREE REINFORCEMENT LEARNING

436 A significant trend in recent research is Label-Free Reinforcement Learning, which aims to improve
 437 LLM reasoning without human-annotated data. Many such methods use the model’s own outputs as
 438 a reward signal. This includes leveraging sequence-level confidence (Li et al., 2025a; Prabhudesai
 439 et al., 2025), the consistency of answers derived from varied reasoning paths (Zhang et al., 2025a;
 440 Zuo et al., 2025; Zhang et al., 2025b; Zhou et al., 2025b; Prasad et al., 2024), minimizing the output
 441 entropy (Agarwal et al., 2025; Cheng et al., 2025), or even random (Shao et al., 2025) or negative
 442 reward (Zhu et al., 2025). These signals are often used within self-training loops where models
 443 fine-tune on their own most plausible solutions (Shafayat et al., 2025; Zhao et al., 2025b). While
 444 these methods all rely on a pre-existing set of unlabeled problems, *R-Zero* removes the need for any
 445 seed dataset.

446 5.2 SELF-PLAY IN LARGE LANGUAGE MODELS

448 The paradigm of self-play, where models take on dual roles to create a self-improvement loop (Chen
 449 et al., 2024; Zhang et al., 2024), has recently been adapted to improve language models without
 450 human data. This approach has been particularly fruitful in verifiable domains like code generation,
 451 where a “Coder” agent’s program is verified by a “Tester” agent’s unit tests (Lin et al., 2025; Wang
 452 et al., 2025a; Pourcel et al., 2025; Jiang et al., 2025). More advanced frameworks push autonomy
 453 further by learning to generate the problems themselves, creating an adaptive curriculum from a
 454 small seed of examples or from scratch (Zhao et al., 2025a; Li et al., 2025c; Zhou et al., 2025a; Fang
 455 et al., 2025). Our work distinguishes itself by extending this paradigm to general reasoning domains
 456 that lack such verifiable environments, like coding tasks.

457 5.3 REINFORCEMENT LEARNING WITH VERIFIABLE REWARDS (RLVR)

459 Reinforcement Learning with Verifiable Rewards has been widely adopted as a versatile paradigm
 460 for enhancing LLMs across a multitude of tasks (Li et al., 2025d; DeepSeek-AI et al., 2025; Shao
 461 et al., 2024). Its effectiveness is demonstrated in diverse applications such as relation extraction (Dai
 462 et al., 2025), interactive GUI navigation (Shi et al., 2025b) and search-engine utilization (Jin et al.,
 463 2025). While early implementations relied on rule-based verifiers, recent work has begun to explore
 464 more sophisticated, model-based verifiers (Ma et al., 2025; Li et al., 2025b; 2024).

466 6 LIMITATION

468 While *R-Zero* demonstrates strong improvements in reasoning performance across multiple do-
 469 mains, several limitations remain. First, the core mechanism of *R-Zero* relies on domains where
 470 correctness can be objectively verified. The Challenger–Solver co-evolution process depends on
 471 clear, deterministic evaluation signals to produce reliable training feedback. Consequently, applying
 472 *R-Zero* to open-ended or subjective tasks, such as creative writing, dialogue, or preference-driven
 473 generation, remains difficult, as these tasks lack unambiguous correctness criteria. In addition, *R-*
 474 *Zero* currently employs specific labeling and verification strategies that may not generalize to all task
 475 types. Developing more robust and broadly applicable labeling mechanisms would further expand
 476 the range of domains where *R-Zero* can be effectively applied.

478 7 CONCLUSION AND FUTURE WORK

480 In this paper, we introduced *R-Zero*, a fully autonomous self-evolving framework that overcomes
 481 data dependency by having a Challenger and Solver co-evolve to create a self-generating curricu-
 482 lum. Our experiments demonstrate that *R-Zero* significantly improves LLM’s reasoning capability
 483 in multiple domains. Future work could further focus on improving efficiency, exploring more ro-
 484 bust labeling techniques, and expanding *R-Zero* to new domains. Extending this self-evolutionary
 485 paradigm to open-ended generative tasks, such as creative writing or dialogue, where evaluation is
 subjective, remains a significant hurdle for future research.

486 REPRODUCIBILITY STATEMENT
 487

488 To ensure the reproducibility of our research, we provide detailed information regarding our data
 489 and experimental setup. All datasets used in this work are publicly available; we provide details on
 490 data sources, any postprocess steps in Appendix B. A comprehensive list of all hyperparameters for
 491 each experiment can be found in a table in Appendix B and Sec. 3.

492
 493 REFERENCES
 494

495 Shivam Agarwal, Zimin Zhang, Lifan Yuan, Jiawei Han, and Hao Peng. The unreasonable effec-
 496 tiveness of entropy minimization in llm reasoning. *ArXiv preprint*, abs/2505.15134, 2025.

497 Sanghwan Bae, Jiwoo Hong, Min Young Lee, Hanbyul Kim, JeongYeon Nam, et al. Online difficulty
 498 filtering for reasoning oriented reinforcement learning. *ArXiv preprint*, abs/2504.03380, 2025.

499
 500 Akhiad Bercovich, Itay Levy, Izik Golan, Mohammad Dabbah, Ran El-Yaniv, et al. Llama-
 501 nemotron: Efficient reasoning models. *ArXiv preprint*, abs/2505.00949, 2025.

502 Martin Briesch, Dominik Sobania, and Franz Rothlauf. Large language models suffer from their
 503 own output: An analysis of the self-consuming training loop. *ArXiv preprint*, abs/2311.16822,
 504 2023.

505
 506 Zixiang Chen, Yihe Deng, Huizhuo Yuan, Kaixuan Ji, and Quanquan Gu. Self-play fine-tuning
 507 converts weak language models to strong language models. *arXiv preprint arXiv:2401.01335*,
 508 2024.

509
 510 Daixuan Cheng, Shaohan Huang, Xuekai Zhu, Bo Dai, Wayne Xin Zhao, et al. Reasoning with
 511 exploration: An entropy perspective. *ArXiv preprint*, abs/2506.14758, 2025.

512 Karl Cobbe, Vineet Kosaraju, Mo Bavarian, Mark Chen, Heewoo Jun, et al. Training verifiers to
 513 solve math word problems. *ArXiv preprint*, abs/2110.14168, 2021.

514
 515 Runpeng Dai, Tong Zheng, Run Yang, and Hongtu Zhu. R1-re: Cross-domain relationship extraction
 516 with rlrv. *ArXiv preprint*, abs/2507.04642, 2025.

517 DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Jun-Mei Song, et al. Deepseek-r1: Incen-
 518 tivizing reasoning capability in llms via reinforcement learning. *ArXiv preprint*, abs/2501.12948,
 519 2025.

520
 521 Elvis Dohmatob, Yunzhen Feng, Arjun Subramonian, and Julia Kempe. Strong model collapse.
 522 *ArXiv preprint*, abs/2410.04840, 2024a.

523
 524 Elvis Dohmatob, Yunzhen Feng, Pu Yang, François Charton, and Julia Kempe. A tale of tails:
 525 Model collapse as a change of scaling laws. In *Forty-first International Conference on Machine
 526 Learning, ICML 2024, Vienna, Austria, July 21-27, 2024*, 2024b.

527
 528 Xinrun Du, Yifan Yao, Kaijing Ma, Bingli Wang, Tianyu Zheng, et al. Supergpqa: Scaling llm
 529 evaluation across 285 graduate disciplines. *ArXiv preprint*, abs/2502.14739, 2025.

530
 531 Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, et al.
 532 The llama 3 herd of models. *ArXiv preprint*, abs/2407.21783, 2024.

533
 534 Wenkai Fang, Shunyu Liu, Yang Zhou, Kongcheng Zhang, Tongya Zheng, et al. Serl: Self-play rein-
 535 forcement learning for large language models with limited data. *ArXiv preprint*, abs/2505.20347,
 536 2025.

537
 538 Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, et al. Olympiadbench:
 539 A challenging benchmark for promoting agi with olympiad-level bilingual multimodal scientific
 540 problems. In *Annual Meeting of the Association for Computational Linguistics*, 2024.

541
 542 Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob
 543 Steinhardt. Measuring massive multitask language understanding. In *Proc. of ICLR*, 2021a.

- 540 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, et al. Measuring mathematical problem solving with the math dataset. *ArXiv preprint*, abs/2103.03874, 2021b.
 541
 542
- 543 Maggie Huan, Yuetai Li, Tuney Zheng, Xiaoyu Xu, Seungone Kim, et al. Does math reasoning
 544 improve general llm capabilities? understanding transferability of llm reasoning. 2025.
 545
 546
- 547 Chengsong Huang, Langlin Huang, Jixuan Leng, Jiacheng Liu, and Jiaxin Huang. Efficient test-time
 548 scaling via self-calibration. *ArXiv preprint*, abs/2503.00031, 2025.
 549
 550
- 551 Hao Jiang, Qi Liu, Rui Li, Yuze Zhao, Yixiao Ma, Shengyu Ye, Junyu Lu, and Yu Su. Verse:
 552 Verification-based self-play for code instructions. In *Proceedings of the AAAI Conference on
 553 Artificial Intelligence*, volume 39, pp. 24276–24284, 2025.
 554
- 555 Bowen Jin, Hansi Zeng, Zhenrui Yue, Dong Wang, Hamed Zamani, and Jiawei Han. Search-r1:
 556 Training llms to reason and leverage search engines with reinforcement learning. *ArXiv preprint*,
 557 abs/2503.09516, 2025.
 558
- 559 Nathan Lambert, Jacob Daniel Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, et al.
 560 Tülu 3: Pushing frontiers in open language model post-training. *ArXiv preprint*, abs/2411.15124,
 561 2024.
 562
- 563 Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay V.
 564 Ramasesh, Ambrose Sloane, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam
 565 Neyshabur, Guy Gur-Ari, and Vedant Misra. Solving quantitative reasoning problems with lan-
 566 guage models. In Sanmi Koyejo, S. Mohamed, A. Agarwal, Danielle Belgrave, K. Cho, and
 567 A. Oh (eds.), *Advances in Neural Information Processing Systems 35: Annual Conference on
 568 Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November
 569 - December 9, 2022*, 2022.
- 570 Pengyi Li, Matvey Skripkin, Alexander Zubrey, Andrey Kuznetsov, and Ivan V. Oseledets.
 571 Confidence is all you need: Few-shot rl fine-tuning of language models. *ArXiv preprint*,
 572 abs/2506.06395, 2025a.
 573
- 574 Ruosen Li, Ziming Luo, and Xinya Du. Fg-prm: Fine-grained hallucination detection and mitigation
 575 in language model mathematical reasoning. *ArXiv preprint*, abs/2410.06304, 2024.
 576
- 577 Zongxia Li, Yapei Chang, Yuhang Zhou, Xiyang Wu, Zichao Liang, Yoo Yeon Sung, and Jordan Lee
 578 Boyd-Graber. Semantically-aware rewards for open-ended r1 training in free-form generation.
 579 *ArXiv preprint*, abs/2506.15068, 2025b.
 580
- 581 Zongxia Li, Xiyang Wu, Guangyao Shi, Yubin Qin, Hongyang Du, Tianyi Zhou, Dinesh Manocha,
 582 and Jordan Lee Boyd-Graber. Videohallu: Evaluating and mitigating multi-modal hallucinations
 583 on synthetic video understanding. *ArXiv preprint*, abs/2505.01481, 2025c.
 584
- 585 Zongxia Li, Wenhao Yu, Chengsong Huang, Rui Liu, Zhenwen Liang, Fuxiao Liu, Jingxi Che, Dian
 586 Yu, Jordan Boyd-Graber, Haitao Mi, et al. Self-rewarding vision-language model via reasoning
 587 decomposition. *arXiv preprint arXiv:2508.19652*, 2025d.
 588
- 589 Zi Lin, Sheng Shen, Jingbo Shang, Jason Weston, and Yixin Nie. Learning to solve and verify: A
 590 self-play framework for code and test generation. *ArXiv preprint*, abs/2502.14948, 2025.
 591
- 592 Xueguang Ma, Qian Liu, Dongfu Jiang, Ge Zhang, Zejun Ma, et al. General-reasoner: Advancing
 593 llm reasoning across all domains. *ArXiv preprint*, abs/2505.14652, 2025.
 594
- 595 Julien Pourcel, Cédric Colas, and Pierre-Yves Oudeyer. Self-improving language models for evolu-
 596 tionary program synthesis: A case study on arc-agc. 2025.
 597
- 598 Mihir Prabhudesai, Lili Chen, Alex Ippoliti, Katerina Fragkiadaki, Hao Liu, et al. Maximizing
 599 confidence alone improves reasoning. *ArXiv preprint*, abs/2505.22660, 2025.
 600
- 601 Archiki Prasad, Weizhe Yuan, Richard Yuanzhe Pang, Jing Xu, Maryam Fazel-Zarandi, Mohit
 602 Bansal, Sainbayar Sukhbaatar, Jason Weston, and Jane Yu. Self-consistency preference opti-
 603 mization. *arXiv preprint arXiv:2411.04109*, 2024.
 604

- 594 Mohamed El Amine Seddik, Swei-Wen Chen, Soufiane Hayou, Pierre Youssef, and Mérouane Deb-
 595 bah. How bad is training on synthetic data? a statistical analysis of language model collapse.
 596 *ArXiv preprint*, abs/2404.05090, 2024.
- 597
- 598 Sheikh Shafayat, Fahim Tajwar, Ruslan Salakhutdinov, Jeff Schneider, and Andrea Zanette. Can
 599 large reasoning models self-train? *ArXiv preprint*, abs/2505.21444, 2025.
- 600 Rulin Shao, Shuyue Stella Li, Rui Xin, Scott Geng, Yiping Wang, et al. Spurious rewards: Rethink-
 601 ing training signals in rlrv. *ArXiv preprint*, abs/2506.10947, 2025.
- 602
- 603 Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Jun-Mei Song, et al. Deepseekmath: Pushing
 604 the limits of mathematical reasoning in open language models. *ArXiv preprint*, abs/2402.03300,
 605 2024.
- 606
- 607 Taiwei Shi, Yiyang Wu, Linxin Song, Tianyi Zhou, and Jieyu Zhao. Efficient reinforcement finetun-
 608 ing via adaptive curriculum learning. *ArXiv preprint*, abs/2504.05520, 2025a.
- 609
- 610 Yucheng Shi, Wenhao Yu, Zaitang Li, Yonglin Wang, Hongming Zhang, et al. Mobilegui-rl: Ad-
 611 vancing mobile gui agent through reinforcement learning in online environment. 2025b.
- 612
- 613 Mehrangiz shoaa kazemi, Bahare Fatemi, Hritik Bansal, John Palowitch, Chrysovalantis Anastasiou,
 614 et al. Big-bench extra hard. In *Annual Meeting of the Association for Computational Linguistics*,
 2025.
- 615
- 616 Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal.
 617 Ai models collapse when trained on recursively generated data. *Nature*, 631, 2024.
- 618
- 619 Yi Su, Dian Yu, Linfeng Song, Juntao Li, Haitao Mi, et al. Crossing the reward bridge: Expanding
 620 rl with verifiable rewards across diverse domains. *ArXiv preprint*, abs/2503.23829, 2025.
- 621
- 622 Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung,
 623 Aakanksha Chowdhery, Quoc Le, Ed Chi, Denny Zhou, and Jason Wei. Challenging BIG-bench
 624 tasks and whether chain-of-thought can solve them. In Anna Rogers, Jordan Boyd-Graber, and
 Naoaki Okazaki (eds.), *Findings of the Association for Computational Linguistics: ACL 2023*,
 2023.
- 625
- 626 Zhen Tan, Dawei Li, Song Wang, Alimohammad Beigi, Bohan Jiang, Amrita Bhattacharjee, Man-
 627 sooreh Karami, Jundong Li, Lu Cheng, and Huan Liu. Large language models for data annotation
 628 and synthesis: A survey. In *Proc. of EMNLP*, 2024.
- 629
- 630 Zhengwei Tao, Ting-En Lin, Xiancai Chen, Hangyu Li, Yuchuan Wu, et al. A survey on self-
 631 evolution of large language models. *ArXiv preprint*, abs/2404.14387, 2024.
- 632
- 633 Yinjie Wang, Ling Yang, Ye Tian, Ke Shen, and Mengdi Wang. Co-evolving llm coder and unit
 tester via reinforcement learning. *ArXiv preprint*, abs/2506.03136, 2025a.
- 634
- 635 Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming
 636 Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi
 637 Fan, Xiang Yue, and Wenhui Chen. Mmlu-pro: A more robust and challenging multi-task language
 638 understanding benchmark. In Amir Globersons, Lester Mackey, Danielle Belgrave, Angela Fan,
 Ulrich Paquet, Jakub M. Tomczak, and Cheng Zhang (eds.), *Advances in Neural Information
 Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024*,
 639 *NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*, 2024.
- 640
- 641 Zengzhi Wang, Fan Zhou, Xuefeng Li, and Pengfei Liu. Octothinker: Mid-training incentivizes
 642 reinforcement learning scaling. *ArXiv preprint*, abs/2506.20512, 2025b.
- 643
- 644 An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, et al. Qwen3 technical report.
 645 *ArXiv preprint*, abs/2505.09388, 2025.
- 646
- 647 Jixiao Zhang and Chunsheng Zuo. Grpo-lead: A difficulty-aware reinforcement learning approach
 for concise mathematical reasoning in language models. *ArXiv preprint*, abs/2504.09696, 2025.

- 648 Kongcheng Zhang, Qi Yao, Shunyu Liu, Yingjie Wang, Baisheng Lai, et al. Consistent paths lead to
 649 truth: Self-rewarding reinforcement learning for llm reasoning. *ArXiv preprint*, abs/2506.08745,
 650 2025a.
- 651
- 652 Qingyang Zhang, Haitao Wu, Changqing Zhang, Peilin Zhao, and Yatao Bian. Right question
 653 is already half the answer: Fully unsupervised llm reasoning incentivization. *ArXiv preprint*,
 654 abs/2504.05812, 2025b.
- 655
- 656 Ruize Zhang, Zelai Xu, Chengdong Ma, Chao Yu, Wei-Wei Tu, Wenhao Tang, Shiyu Huang, Deheng
 657 Ye, Wenbo Ding, Yaodong Yang, et al. A survey on self-play methods in reinforcement learning.
 658 *arXiv preprint arXiv:2408.01072*, 2024.
- 659
- 660 Andrew Zhao, Yiran Wu, Yang Yue, Tong Wu, Quentin Xu, et al. Absolute zero: Reinforced self-
 661 play reasoning with zero data. *ArXiv preprint*, abs/2505.03335, 2025a.
- 662
- 663 Xuandong Zhao, Zhewei Kang, Aosong Feng, Sergey Levine, and Dawn Xiaodong Song. Learning
 664 to reason without external rewards. *ArXiv preprint*, abs/2505.19590, 2025b.
- 665
- 666 Yulai Zhao, Haolin Liu, Dian Yu, S. Y. Kung, Haitao Mi, and Dong Yu. One token to fool llm-as-a-
 667 judge. volume abs/2507.08794, 2025c.
- 668
- 669 Tong Zheng, Lichang Chen, Simeng Han, R Thomas McCoy, and Heng Huang. Learning to reason
 670 via mixture-of-thought for logical reasoning. *arXiv preprint arXiv:2505.15817*, 2025a.
- 671
- 672 Yaowei Zheng, Junting Lu, Shenzhi Wang, Zhangchi Feng, Dongdong Kuang, et al. Easyr1: An
 673 efficient, scalable, multi-modality rl training framework. 2025b.
- 674
- 675 Yifei Zhou, Sergey Levine, Jason E. Weston, Xian Li, and Sainbayar Sukhbaatar. Self-challenging
 676 language model agents. *ArXiv preprint*, abs/2506.01716, 2025a.
- 677
- 678 Yujun Zhou, Zhenwen Liang, Haolin Liu, Wenhao Yu, Kishan Panaganti, Linfeng Song, Dian Yu,
 679 Xiangliang Zhang, Haitao Mi, and Dong Yu. Evolving language models without labels: Majority
 680 drives selection, novelty promotes variation. *arXiv preprint arXiv:2509.15194*, 2025b.
- 681
- 682 Yujun Zhou, Jiayi Ye, Zipeng Ling, Yufei Han, Yue Huang, Haomin Zhuang, Zhenwen Liang, Kehan
 683 Guo, Taicheng Guo, Xiangqi Wang, et al. Dissecting logical reasoning in llms: A fine-grained
 684 evaluation and supervision study. *arXiv preprint arXiv:2506.04810*, 2025c.
- 685
- 686 Xinyu Zhu, Mengzhou Xia, Zhepei Wei, Wei-Lin Chen, Danqi Chen, et al. The surprising effectiveness
 687 of negative reinforcement in llm reasoning. *ArXiv preprint*, abs/2506.01347, 2025.
- 688
- 689 Yuxin Zuo, Kaiyan Zhang, Shang Qu, Li Sheng, Xuekai Zhu, et al. Ttrl: Test-time reinforcement
 690 learning. *ArXiv preprint*, abs/2504.16084, 2025.
- 691
- 692
- 693
- 694
- 695
- 696
- 697
- 698
- 699
- 700
- 701

APPENDIX

A THE USE OF LARGE LANGUAGE MODELS (LLMs)

We acknowledge the use of large language models (LLMs) as assistive tools in this research, with usage limited to refining grammar and improving language clarity in the manuscript, writing utility scripts for data preprocessing and postprocessing, and debugging; all outputs from these models were meticulously reviewed, revised, and verified by the authors, who retain full responsibility for all content presented in this paper.

B EXPERIMENT DETAILS

B.1 TRAINING HYPERPARAMETER

This section summarizes the most critical algorithmic hyperparameters for the Solver and Challenger training stages. All experiments were conducted using BFloat16 (BF16) mixed precision and FlashAttention 2.

B.1.1 SOLVER TRAINING

- **Global Batch Size:** 128
- **Learning Rate:** 1×10^{-6}
- **Weight Decay:** 1×10^{-2}
- **KL Penalty Coefficient** (λ_{KL}): 1×10^{-2}
- **Max Steps:** 15
- **Number of Rollouts:** 5
- **Rollout Temperature:** 1.0
- **Rollout Top-p:** 0.99

B.1.2 CHALLENGER TRAINING

- **Global Batch Size:** 128
- **Learning Rate:** 1×10^{-6}
- **Weight Decay:** 1×10^{-2}
- **KL Penalty Coefficient** (λ_{KL}): 1×10^{-2}
- **Max Steps:** 5
- **Number of Rollouts:** 4
- **Rollout Temperature:** 1.0
- **Rollout Top-p:** 0.99

B.2 PROMPT TEMPLATES

This section presents the exact prompt templates used for the solver and challenger models.

Solver Prompt Template

System Message:

Please reason step by step, and put your final answer within `\boxed{}`.

User Message:

`{problem_statement}`

Note: {problem_statement} is a placeholder for the actual math problem.

756 **Challenger Prompt Template**
 757
 758 **System Message:**
 759 You are an expert competition-math problem setter. FIRST, in your private scratch-pad, think
 760 step-by-step to design a brand-new, non-trivial problem. The problem could come from
 761 any field of mathematics, including but not limited to algebra, geometry, number theory,
 762 combinatorics, prealgebra, probability, statistics, and calculus. Aim for a difficulty such that
 763 fewer than 30% of advanced high-school students could solve it. Avoid re-using textbook
 764 clichés or famous contest problems.
 765 THEN, without revealing any of your private thoughts, output **exactly** the following two
 766 blocks:
 767 <question>
 768 {The full problem statement on one or more lines}
 769 </question>
 770 \boxed{\text{final_answer}}
 771
 772 Do NOT output anything else—no explanations, no extra markup.
 773 **User Message:**
 774 Generate one new, challenging reasoning question now. Remember to format the output
 775 exactly as instructed.

778 B.3 GPT-4O JUDGE PROMPT

779
 780 To programmatically evaluate the correctness of answers on mathematical benchmarks where the
 781 final answer can be complex (e.g., simplified expressions), we use GPT-4o as a judge. The exact
 782 prompt and configuration used for this evaluation are detailed below.

783 Configuration for GPT-4o as Judge

- 785 • **Model:** gpt-4o
- 786 • **Temperature:** 0.1

788 **System Message:**

789 You are a math answer checker.

790 **User Message Template:**

791 Hi, there is an answer: {answer},
 792 and the ground truth answer is: {response},
 793 please check whether the answer is correct or not, and return the ****only****
 794 Yes or No.

795 *Note: {answer} is a placeholder for the model-generated solution, and {response} is
 796 the ground-truth answer from the benchmark.*

800 B.4 REPETITION PENALTY IMPLEMENTATION

801
 802 To encourage the Challenger to generate a diverse set of questions within each batch, we apply a
 803 repetition penalty, r_{rep} . This penalty is designed to disincentivize the model from producing seman-
 804 tically similar questions in the same batch. The implementation is a multi-step process based on
 805 clustering questions by their BLEU score similarity.

806 **1. Pairwise Distance Calculation via BLEU Score** First, we compute a pairwise distance matrix
 807 for all questions in a batch. The distance d_{ij} between any two questions, x_i and x_j , is defined as
 808 one minus their BLEU score:
 809

$$d_{ij} = 1 - \text{BLEU}(x_i, x_j)$$

For this calculation, we specifically use the `sentence_bleu` function from the NLTK library (`nltk.translate.bleu_score`). To ensure numerical stability, especially for shorter questions with limited n-gram overlap, we employ its first smoothing function, `SmoothingFunction().method1`. The questions are tokenized for the BLEU calculation by splitting on whitespace; no further text normalization, such as lowercasing or punctuation removal, is performed.

2. Agglomerative Clustering With the pairwise distance matrix computed, we then group similar questions using agglomerative hierarchical clustering. This step is performed using the `Clustering` implementation from the `scikit-learn` library. The clustering algorithm is configured with the following key parameters:

- **Metric:** Set to ‘`precomputed`’, indicating that we provide our custom BLEU-based distance matrix instead of having the algorithm compute distances.
- **Linkage:** Set to ‘`average`’. This method defines the distance between two clusters as the average of the distances between all pairs of questions across the two clusters.

3. Final Penalty Calculation Once each question in the batch is assigned to a cluster, the repetition penalty $r_{\text{rep}}(x_i)$ for a given question x_i is determined by the relative size of the cluster C_k to which it belongs. The penalty is calculated as:

$$r_{\text{rep}}(x_i) = \frac{|C_k|}{B}$$

Here, $|C_k|$ represents the number of questions in cluster C_k , and B is the total number of questions in the batch (i.e., the batch size).

C EVALUATION BENCHMARK

We assess our framework on a comprehensive suite of benchmarks. Although the question-generator prompt for our method is primarily focused on mathematical problem-solving, a key objective of our evaluation is to explore whether the resulting improvements in reasoning ability can generalize to other domains. Therefore, our evaluation is divided into two main categories.

Mathematical Reasoning. We use seven challenging benchmarks: AMC, Minerva (Lewkowycz et al., 2022), MATH-500 (Hendrycks et al., 2021b), GSM8K (Cobbe et al., 2021), Olympiad-Bench (He et al., 2024), AIME-2024, and AIME-2025. For these tasks, where answers can be complex, we employ GPT-4o as a programmatic judge to semantically verify the correctness of the final answer against the ground truth (Zhao et al., 2025c). For the difficult AMC and AIME benchmarks, we report the `mean@32` metric. For all other math benchmarks, we report accuracy based on greedy decoding.

General Domain Reasoning. To test for the generalization of reasoning ability, we evaluate the following challenging benchmarks:

- **MMLU-Pro** (Wang et al., 2024): An enhanced version of the MMLU (Hendrycks et al., 2021a) benchmark, featuring a more challenging suite of multi-task questions designed to provide a stricter evaluation of language model capabilities.
- **SuperGPQA** (Du et al., 2025): A large-scale benchmark focused on graduate-level reasoning. It comprises questions across 285 distinct disciplines that have been verified as unsearchable on the web, thereby isolating true reasoning ability from simple knowledge recall.
- **BBEH** (shoaa kazemi et al., 2025): This benchmark builds upon the foundation of BIG-Bench Hard (Suzgun et al., 2023) by incorporating a new selection of tasks specifically engineered to be more difficult, thus providing a more accurate measure of complex reasoning skills.

For this category, we follow the experimental setup, prompts, and evaluation codes from (Ma et al., 2025), reporting Exact Match (EM) accuracy obtained via greedy decoding.

864 D PARAMETER SHARING BETWEEN CHALLENGER AND SOLVER
865

866  Separating the Challenger and Solver into two models will be better.
867

868
869 Table 6: Comparison of math performance and pseudo-label accuracy between the standard R-Zero
870 (two-model) and Single-R-Zero (unified model, shared parameters) frameworks across iterations.
871

872

Iteration	R-Zero (ours)		Single-R-Zero	
	Performance	Pseudo-label Acc (%)	Performance	Pseudo-label Acc (%)
Step 15	48.06	71.0	47.31	63.4
Step 30	48.44	56.2	46.95	46.6
Step 45	49.12	48.8	45.57	32.6
Step 60	46.52	42.2	43.89	33.8

873
874 To investigate whether the separation of the Challenger and Solver into two independent models
875 is a necessary component for the success of *R-Zero*, we conduct an ablation study using a unified
876 model with shared parameters. In this configuration (Single-R-Zero), a single model is tasked with
877 performing both roles, i.e., generating a challenging curriculum and subsequently learning from it.
878

879 The results, presented in Table 6, clearly indicate that separating the Challenger and Solver into two
880 independent models is crucial for both performance and stability. We observe two key findings. First,
881 our standard two-model R-Zero framework not only achieves a higher peak performance (49.12) but
882 also sustains improvement for more iterations, with its collapse occurring after the third iteration.
883 In contrast, the unified Single-R-Zero model’s performance peaks after the very first iteration and
884 degrades immediately thereafter. Second, the Single-R-Zero model, where the agent must generate
885 and solve its own problems, produces pseudo-labels of significantly lower accuracy at every stage.
886 For example, in the first iteration, its pseudo-label accuracy is already substantially lower than the
887 R-Zero’s (63.4% vs. 71.0%). We hypothesize that this is because having the problem-setter and
888 solver originate from the same model leads to a form of overconfidence that comes from internal
889 bias.
890

891 E BEYOND LABEL NOISE: UNPACKING THE ROOTS OF INSTABILITY
892

893 The most immediate hypothesis for this performance col-
894 lapse is the degradation of pseudo-label quality, a pot-
895 tential failure mode of the self-correction mechanism we
896 discussed in Section 4.4. As the Challenger generates
897 increasingly difficult problems, it is plausible that the
898 Solver’s majority vote becomes a less reliable source for
899 ground truth, resulting in a noisy training signal that could
900 ultimately harm performance. To empirically test the
901 extent to which this is the primary cause, we sampled
902 500 questions from a later training iteration to conduct a
903 more granular investigation into the relationship between
904 pseudo-label fidelity and the observed performance drop.
905

906 Table 7: Accuracy of self-generated
907 pseudo-labels (%), labeled by Gemini.
908 Shaded and bolded values indicate the
909 best checkpoint for each model size.
910

911

Step	Model Size		
	0.6B	1.7B	4B
Step 15	70.6	69.4	71.0
Step 30	53.4	55.2	56.2
Step 45	50.8	52.2	48.8
Step 60	44.0	45.2	42.2

912 Although the degradation of pseudo-label accuracy is a consistent trend across iterations, our anal-
913 ysis suggests this is not the primary, nor even the sole, driver of the eventual performance collapse.
914 Table 7 presents the pseudo-label data quality for each model at the onset of its performance col-
915 lapse. Intriguingly, there appears to be no universal accuracy threshold that triggers this degradation.
916 For instance, the 0.6B model begins to decline when data accuracy is still as high as 70.6% (Step
917 15), whereas the 4B model tolerates an accuracy as low as 48.8% (Step 45) before its performance
918 drops.
919

920 This suggests that the absolute percentage of label noise is not the sole determinant of instability.
921 Another potential, and perhaps more fundamental, reason is a form of model collapse that can be
922 introduced when training exclusively on self-synthesized data (Tan et al., 2024; Shumailov et al.,
923 2024; Dohmatob et al., 2024b; Zhou et al., 2025c; Seddik et al., 2024; Dohmatob et al., 2024a;
924

918 Briesch et al., 2023; Zheng et al., 2025a). A model can enter a degenerative feedback loop, suffering
 919 from a loss of diversity or an amplification of its own biases, which presents a significant challenge.
 920

921 F THEORETICAL ANALYSIS

923 In this section, we provide a theoretical motivation for our uncertainty reward function, $r_{\text{uncertainty}} \propto$
 924 $1 - 2|\hat{p}(x; S_\phi) - \frac{1}{2}|$, which is maximized when the Solver’s success probability, \hat{p} , is 50%. Our
 925 analysis is grounded in recent work that formally establishes that the most efficient training occurs
 926 when a learner is exposed to tasks at the frontier of its capabilities (Shi et al., 2025a; Bae et al.,
 927 2025).

928 The core insight from these studies is that the learning potential of the current Solver, with policy
 929 S_ϕ , can be quantified by the KL divergence to an optimal policy S^* . This divergence, $\mathbb{D}_{KL}(S_\phi || S^*)$,
 930 is lower-bounded by the variance of the Solver’s reward. For the binary reward signal used in our
 931 framework, the success probability is \hat{p} . This leads to the specific lower bound:

$$933 \quad \mathbb{D}_{KL}(S_\phi || S^*) \geq \frac{\hat{p}(1 - \hat{p})}{2\beta^2}$$

935 where β is the temperature parameter that controls entropy regularization. The right-hand side of
 936 the inequality, which is proportional to the reward variance, is maximized precisely when $\hat{p} = 0.5$.
 937 Therefore, by designing the Challenger’s reward to incentivize questions that push the current Solver
 938 towards this point of maximum uncertainty, our framework is theoretically motivated to generate a
 939 maximally efficient curriculum in each iteration of the co-evolutionary process.

972 **G ALGORITHM OF R-Zero**
973

974 Here, we provide the pseudocode for our method.
975

976 **Algorithm 1: R-Zero: Self-Evolving Challenger-Solver Framework**
977

978 **Input:** Initial models Q_θ, S_ϕ ; Group size G ; Solver samples m ; Dataset size N ; Filtering
979 threshold δ ; Repetition threshold τ_{BLEU} . **Output:** Evolved models Q_θ and S_ϕ .
980 **for each self-play iteration do**

981 // --- Phase 1: Challenger Training (Sec 2.2) ---
982 **for each Challenger training step do**
983 Sample question group $\{x_i\}_{i=1}^G \sim Q_\theta(\cdot)$;
984 **for each question x_i do**
985 **if** $\text{FormatCheck}(x_i)$ is invalid (e.g., missing tags) **then**
986 | $r_i \leftarrow 0$;
987 **end**
988 **else**
989 Sample m answers $\{y_j\}_{j=1}^m \sim S_\phi(\cdot | x_i)$;
990 Get pseudo-label $\tilde{y}_i \leftarrow \text{MajorityVote}(\{y_j\}_{j=1}^m)$;
991 Compute correctness $\hat{p}_i \leftarrow \frac{1}{m} \sum_{j=1}^m \mathbb{1}\{y_j = \tilde{y}_i\}$;
992 Compute uncertainty reward $r_{\text{uncertainty}} \leftarrow 1 - 2|\hat{p}_i - \frac{1}{2}|$;
993 Compute repetition penalty $r_{\text{rep}}(x_i)$ via BLEU clustering (Sec 2.2);
994 Final reward: $r_i \leftarrow \max(0, r_{\text{uncertainty}} - r_{\text{rep}}(x_i))$;
995 **end**
996 **end**
997 Update Q_θ via GRPO using rewards $\{r_i\}_{i=1}^G$;
998 **end**
999 // --- Phase 2: Solver Dataset Construction (Sec 2.3) ---
1000 Initialize curated dataset $\mathcal{S} \leftarrow \emptyset$;
1001 Sample N candidate questions $\{x_k\}_{k=1}^N \sim Q_\theta(\cdot)$;
1002 **for each candidate x_k do**
1003 Sample m answers $\{y_j\}_{j=1}^m \sim S_\phi(\cdot | x_k)$;
1004 Get pseudo-label $\tilde{y}_k \leftarrow \text{MajorityVote}(\{y_j\})$;
1005 Compute correctness $\hat{p}_k \leftarrow \frac{1}{m} \sum_{j=1}^m \mathbb{1}\{y_j = \tilde{y}_k\}$;
1006 **if** $|\hat{p}_k - \frac{1}{2}| \leq \delta$ **then**
1007 | Add (x_k, \tilde{y}_k) to \mathcal{S} ;
1008 **end**
1009 **end**
1010 // --- Phase 3: Solver Training (Sec 2.4) ---
1011 **for each minibatch $(x, \tilde{y}) \in \mathcal{S}$ do**
1012 Sample G answers $\{y_j\}_{j=1}^G \sim S_\phi(\cdot | x)$;
1013 Compute binary rewards $r_j \leftarrow \mathbb{1}(y_j = \tilde{y})$;
1014 Update S_ϕ via GRPO using rewards $\{r_j\}_{j=1}^G$;
1015 **end**
1016 **end**

1017 **H GENERALIZING BEYOND THE MATH DOMAIN**
1018

1019 To further evaluate the generality of our proposed approach, we conduct an additional experiment in
1020 which the model is no longer instructed to generate exclusively mathematical problems during data
1021 construction. Instead, we remove the math-specific constraint from the prompting stage, allowing
1022 the model to produce a broader mixture of questions, including commonsense, logical reasoning, and
1023 other non-mathematical categories. This setup enables us to examine whether the improvements ob-
1024 served in earlier experiments are merely a consequence of alignment with math-heavy benchmarks
1025 or whether our method can maintain (or even enhance) performance under a more domain-general
1026 training regime.

1026 Table 8: Results of different models and prompting strategies on mathematics and reasoning benchmarks.
 1027
 1028

	AMC	Minerva	MATH	GSM8K	Olympiad	AIME25	AIME24	SuperGPQA	MMLU-Pro	BBEH
Qwen3-4B	45.7	38.24	68.2	87.79	41.04	10.3	6.7	20.88	37.38	7.57
R-Zero	57.27	52.94	79.6	92.12	44.59	9.6	13.4	27.55	55.47	10.42
R-Zero (no math)	56.27	52.56	79.2	92.56	43.56	9.9	13.2	28.14	54.75	10.31
Qwen3-8B	51.95	50	78	89.08	44.74	12.1	14.6	28.33	51.8	8.63
R-Zero	61.67	60.66	82	94.09	48.89	13.3	15.4	31.38	61.53	10.60
R-Zero (no math)	65.53	61.19	81.6	93.89	49.24	13.2	15.2	32.29	61.83	10.88

1034
 1035
 1036 The motivation behind introducing this experiment stems from concerns that many mathematical
 1037 benchmarks contain well-structured and heavily studied problem formats. If the training data is re-
 1038 stricted to similar mathematical tasks, it may artificially inflate performance gains while providing
 1039 limited evidence regarding the method’s real-world applicability. By diversifying the data genera-
 1040 tion process, we aim to assess whether the model can benefit from richer and more heterogeneous
 1041 supervision signals.

1042 The experimental results, summarized in Table 8, demonstrate that when the base model has suffi-
 1043 cient capacity (e.g., the 8B variant), training on more general data leads to further improvements in
 1044 both math and non-math benchmarks. Notably, the “R-Zero (remove math prompt)” setting exhibits
 1045 the highest performance on the general benchmarks for the 8B model, suggesting that our approach
 1046 effectively transfers to broader reasoning tasks beyond pure mathematics.

1047 These findings provide additional evidence that the proposed method is not limited to mathematical
 1048 domains; rather, it scales effectively with model capacity and supports improved reasoning perfor-
 1049 mance under more diverse training conditions.

I PRELIMINARIES

1053 Our work builds upon recent advancements in reinforcement learning for fine-tuning large language
 1054 models. We briefly review two key methodologies that are relevant to our framework.

I.1 GROUP RELATIVE POLICY OPTIMIZATION

1055 Group Relative Policy Optimization (GRPO) (Shao et al., 2024) is a reinforcement learning algo-
 1056 rithm for fine-tuning policy LLMs π_θ without a separately learned value function. Its key idea is to
 1057 normalize rewards within a group of responses sampled from the same prompt, thereby stabilizing
 1058 optimization.

1059 **Setup.** Given a query \mathbf{q} , the old policy $\pi_{\theta_{\text{old}}}$ generates G candidate responses $\{\mathbf{o}_1, \dots, \mathbf{o}_G\}$. Each
 1060 response \mathbf{o}_i is assigned a scalar reward $R(\mathbf{q}, \mathbf{o}_i)$. Group-wise z-score normalization is then applied
 1061 to obtain an advantage shared across the tokens of the response:

$$\hat{A}_{i,t} = \frac{R(\mathbf{q}, \mathbf{o}_i) - \text{mean}(\{R(\mathbf{q}, \mathbf{o}_1), \dots, R(\mathbf{q}, \mathbf{o}_G)\})}{\text{std}(\{R(\mathbf{q}, \mathbf{o}_1), \dots, R(\mathbf{q}, \mathbf{o}_G)\}) + \varepsilon_{\text{norm}}},$$

1062 where $\varepsilon_{\text{norm}}$ is a small constant for numerical stability. The same normalized advantage $\hat{A}_{i,t}$ is
 1063 applied to all tokens of response \mathbf{o}_i .

1064 **Policy Update.** Let $r_\theta(o_{i,t}) = \frac{\pi_\theta(o_{i,t} | \mathbf{q}, \mathbf{o}_{i,<t})}{\pi_{\theta_{\text{old}}}(o_{i,t} | \mathbf{q}, \mathbf{o}_{i,<t})}$. The policy is updated with a clipped surrogate
 1065 objective, similar to PPO, combined with a KL-divergence penalty to constrain policy drift:

$$\mathcal{L}_{\text{GRPO}}(\theta) = -\frac{1}{G} \sum_{i=1}^G \frac{1}{|\mathbf{o}_i|} \sum_{t=1}^{|\mathbf{o}_i|} \min(r_\theta(o_{i,t}) \hat{A}_{i,t}, \text{clip}(r_\theta(o_{i,t}), 1 - \epsilon, 1 + \epsilon) \hat{A}_{i,t}) + \beta \text{KL}(\pi_\theta(\mathbf{q}) \| \pi_{\theta_{\text{old}}}(\mathbf{q}))$$

1066 Maximizing the negative of this loss encourages the policy to increase the probability of tokens
 1067 contributing to responses with positive relative advantages, while clipping prevents overly aggressive
 1068 updates. The KL penalty, weighted by β , further stabilizes training by preventing the new policy
 1069 from drifting too far from the old one.

1080
1081

I.2 REINFORCEMENT LEARNING WITH VERIFIABLE REWARDS

1082
1083
1084

Reinforcement Learning with Verifiable Rewards (RLVR) (Lambert et al., 2024) is a paradigm for fine-tuning models in domains where response quality can be deterministically verified. RLVR relies on a rule-based verifier $v : \mathcal{X} \rightarrow \{0, 1\}$ that assigns a binary reward to each response x_i :

1085
1086
1087

$$r_i = v(x_i) = \begin{cases} 1, & \text{if } x_i \text{ satisfies a task-specific correctness check,} \\ 0, & \text{otherwise.} \end{cases}$$

1088
1089
1090

This reward structure is especially effective for tasks like math, code generation with clear correctness criteria, and serves as the foundation for the reward mechanism in our Solver training.

1091
1092
1093
1094
1095
1096
1097
1098
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133